

```
In [42]: import pandas as pd
import numpy as np
import xgboost as xgb
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestRegressor

In [43]: df = pd.read_csv("vegiedata.csv")
df.head(10)

Out [43]:
```

	Sl no.	District Name	Market Name	Commodity	Variety	Grade	Min Price	Max Price	Modal Price	Date
0	1	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	30-Dec-22
1	2	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	29-Dec-22
2	3	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	28-Dec-22
3	4	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	23-Dec-22
4	5	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	22-Dec-22
5	6	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	21-Dec-22
6	7	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	20-Dec-22
7	8	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	19-Dec-22
8	9	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	17-Dec-22
9	10	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	15-Dec-22

```
In [44]: import pandas as pd

# Create a dictionary to map the old column names to the new ones
column_mapping = {'District Name': 'district', 'Market Name': 'market', 'Date': 'arrival_date',
                  'Max Price': 'max_price', 'Min Price': 'min_price', 'Min Price': 'modal_price',
                  'Commodity': 'commodity', 'Variety': 'variety'}

# Rename the columns using the mapping dictionary
df.rename(columns=column_mapping, inplace=True)

In [45]: data = df

In [ ]:

In [ ]:

In [46]: data

Out [46]:
```

	Sl no.	district	market	commodity	variety	Grade	min_price	max_price	modal_price	arrival_date
0	1	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	30-Dec-22
1	2	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	29-Dec-22
2	3	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	28-Dec-22
3	4	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	23-Dec-22
4	5	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	22-Dec-22
...
5180	5183	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
5181	5184	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
5182	5185	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
5183	5186	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
5184	5187	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

5185 rows × 10 columns

```
In [47]: data = data.dropna()

In [48]: data

Out [48]:
```

	Sl no.	district	market	commodity	variety	Grade	min_price	max_price	modal_price	arrival_date
0	1	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	30-Dec-22
1	2	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	29-Dec-22
2	3	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	28-Dec-22
3	4	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	23-Dec-22
4	5	Adilabad	Adilabad(Rythu Bazar)	Tomato	Other	FAQ	1500.0	1500.0	1500.0	22-Dec-22
...
3993	3996	Mahbubnagar	Kalwakurthy	Brinjal	Brinjal	FAQ	500.0	1500.0	1000.0	20-May-22
3994	3997	Mahbubnagar	Kalwakurthy	Brinjal	Brinjal	FAQ	500.0	1100.0	800.0	19-May-22
3995	3998	Mahbubnagar	Kalwakurthy	Brinjal	Brinjal	FAQ	600.0	1500.0	1050.0	19-May-22
3996	3999	Mahbubnagar	Kalwakurthy	Brinjal	Brinjal	FAQ	1000.0	1500.0	1250.0	17-May-22
3997	4000	Mahbubnagar	Kalwakurthy	Brinjal	Brinjal	FAQ	1000.0	1500.0	1250.0	16-May-22

3998 rows × 10 columns

```
In [49]: data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 3998 entries, 0 to 3997
Data columns (total 10 columns):
 #  Column      Non-Null Count  Dtype
---  --
 0  sl no.      3998 non-null   int64
 1  district    3998 non-null   object
 2  market      3998 non-null   object
 3  commodity    3998 non-null   object
 4  variety      3998 non-null   object
 5  grade       3998 non-null   object
 6  min_price   3998 non-null   float64
 7  max_price   3998 non-null   float64
 8  modal_price 3998 non-null   float64
 9  arrival_date 3998 non-null   object
dtypes: float64(3), int64(1), object(6)
memory usage: 343.6+ KB

In [50]: data['min_price'] = data['min_price']/100
data['max_price'] = data['max_price']/100
data['modal_price'] = data['modal_price'] / 100

C:\Users\koppil_kutdpvm\AppData\Local\Temp\ipykernel_26604\3474568116.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
data['min_price'] = data['min_price']/100
C:\Users\koppil_kutdpvm\AppData\Local\Temp\ipykernel_26604\3474568116.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
data['max_price'] = data['max_price']/100
C:\Users\koppil_kutdpvm\AppData\Local\Temp\ipykernel_26604\3474568116.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
data['modal_price'] = data['modal_price'] / 100

In [51]: X1 = data.iloc[:,8]

In [52]: Y1 = data['modal_price']

In [53]: X = data[['commodity','min_price','max_price']]
Y = data['modal_price']
cat_mask = (X.dtypes==object)
cat_cols = X.columns[cat_mask].tolist()
le = LabelEncoder()
X[cat_cols] = X[cat_cols].apply(lambda x:le.fit_transform(x))

C:\Users\koppil_kutdpvm\AppData\Local\Temp\ipykernel_26604\1597142809.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
X[cat_cols] = X[cat_cols].apply(lambda x:le.fit_transform(x))

In [54]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.2,random_state=42)

In [55]: def xgb_model():
    st = StandardScaler()
    xgb_req = xgb.XGBRegressor()
    steps = [('scaler',st),('model',xgb_req)]
    xgb.pipeline = Pipeline(steps)

    param = {
        'model__subsample': np.arange(0.05,1.05),
        'model__max_depth': np.arange(3,20,1),
        'model__colsample_bytree': np.arange(1,1.05,.05),
        'model__learning_rate': np.arange(0,1,1)
    }
    rand = RandomizedSearchCV(estimator=xgb.pipeline,param_distributions = param,n_iter=3,scoring='neg_mean_squared_error',cv=4)
    rand.fit(X_train,Y_train)
    model = rand.best_estimator_
    return model

In [56]: model = xgb_model()

In [57]: model.fit(X_train,Y_train)

Out [57]: Pipeline(steps=[('scaler', StandardScaler()),
                        ('model',
                          XGBRegressor(base_score=None, booster=None, callbacks=None,
                                         colsample_bylevel=None, colsample_bynode=None,
                                         colsample_bytree=0.9090909090909091,
                                         early_stopping_rounds=None,
                                         enable_categorical=False, eval_metric=None,
                                         feature_types=None, gamma=None, gpu_id=None,
                                         grow_policy=None, importance_type=None,
                                         interaction_constraints=None, learning_rate=0.4,
                                         max_bin=None, max_cat_threshold=None,
                                         max_cat_to_onehot=None, max_delta_step=None,
                                         max_depth=4, max_leaves=None,
                                         min_child_weight=None, missingnan,
                                         monotone_constraints=None, n_estimators=100,
                                         n_jobs=None, num_parallel_tree=None,
                                         predictor=None, random_state=None, ...))])

In [58]: model.score(X_test,Y_test)

Out [58]: 0.985348328371898

In [59]: model1 = xgb_model()

In [60]: model1.fit(X_train,Y_train)

Out [60]: Pipeline(steps=[('scaler', StandardScaler()),
                        ('model',
                          XGBRegressor(base_score=None, booster=None, callbacks=None,
                                         colsample_bylevel=None, colsample_bynode=None,
                                         colsample_bytree=0.9090909090909091,
                                         early_stopping_rounds=None,
                                         enable_categorical=False, eval_metric=None,
                                         feature_types=None, gamma=None, gpu_id=None,
                                         grow_policy=None, importance_type=None,
                                         interaction_constraints=None, learning_rate=0.1,
                                         max_bin=None, max_cat_threshold=None,
                                         max_cat_to_onehot=None, max_delta_step=None,
                                         max_depth=8, max_leaves=None,
                                         min_child_weight=None, missingnan,
                                         monotone_constraints=None, n_estimators=100,
                                         n_jobs=None, num_parallel_tree=None,
                                         predictor=None, random_state=None, ...))])

In [79]: model1.score(X_test,Y_test)

Out [79]: 0.9805496845820512

In [98]: import pickle
import xgboost as xgb

# Save the trained model to a file
#filename = 'xgboost_model.pkl'
#pickle.dump(model, open(filename, 'wb'))

# Load the saved model from the file
#loaded_model = pickle.load(open(filename, 'rb'))

In [86]: import pickle
import xgboost as xgb

# Save the trained model to a file
#filename = 'xgb_model1.sav'
#pickle.dump(model1, open(filename, 'wb'))

In [133]: import numpy as np
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
import xgboost as xgb
from sklearn.tree import DecisionTreeRegressor

# Split the data into train and test sets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

# Encode categorical features
cat_mask = (X_train.dtypes == object)
cat_cols = X_train.columns[cat_mask].tolist()
le = LabelEncoder()
X_train[cat_cols] = X_train[cat_cols].apply(lambda x: le.fit_transform(x))

# Scale the numerical features
st = StandardScaler()
X_train[['min_price', 'max_price']] = st.fit_transform(X_train[['min_price', 'max_price']])

# Define the XGBoost regressor
model = DecisionTreeRegressor()

# Train the XGBoost regressor
model.fit(X_train, Y_train)

# Encode categorical features in the test data
X_test[cat_cols] = X_test[cat_cols].apply(lambda x: le.transform(x))

# Scale the numerical features in the test data
X_test[['min_price', 'max_price']] = st.transform(X_test[['min_price', 'max_price']])

# Make predictions using the trained model
predictions = model.predict(X_test)

# Display the predictions
print("Predicted modal prices:", predictions)

Predicted modal prices: [28.      44.      40.      7.95238095 50.07142857 7.6
 7.26086957 7.60294118 10.42105263 30.      40.
 36.      12.1402439 34.      7.26086957 12.      5.75
 20.      3.5      3.      40.      35.      26.25
 35.      5.86477273 10.      30.      7.60294118 7.26086957
 17.5      13.      5.86477273 12.      15.31147541 22.
 7.26086957 11.      4.18181818 10.42105263 14.      20.
 3.      17.78333333 35.      20.      7.91666667
 8.89473684 13.25      50.07142857 44.      7.26086957 4.65
 6.05405405 5.      20.      44.      7.95238095 15.01408451
 25.02469136 15.      18.5      18.      5.86477273 23.
 10.      22.      13.      12.8      36.
 5.75      40.      50.07142857 15.01408451 50.07142857 25.02469136
 14.66666667 8.07741935 5.97058824 4.83333333 11.      50.07142857
 3.02352941 15.9      44.      27.06666667 40.      15.31147541
 70.      24.      5.60869565 27.06666667 18.
 6.15      13.75      30.      15.33333333 5.60869565 11.5
 7.26086957 13.75      5.      23.      15.      15.31147541
 7.26086957 24.      16.      40.      10.625      12.
 20.      40.      12.1402439 45.      26.5      5.
 25.04347826 3.33333333 15.04347826 50.07142857 12.22222222 21.
 7.95238095 22.75      15.01408451 10.      7.26086957 15.01408451
 18.4      10.      20.      10.      20.      45.
 12.1402439 15.92307692 22.8      15.33333333 7.60294118 17.78333333
 48.      6.05405405 32.      15.31147541 21.      6.05405405
 25.02469136 13.      22.23333333 28.      42.5      6.05405405
 7.26086957 40.16666667 44.      2.      7.91666667 30.
 3.      19.      27.9      15.01408451 30.      50.07142857
 60.      50.      11.25      3.      23.      10.
 8.5      18.      15.31147541 15.01408451 15.25      7.91666667
 25.02469136 8.07741935 10.42105263 7.95238095 28.      24.5
 5.86477273 15.31147541 58.      20.      40.16666667 5.86477273
 30.      7.60294118 15.01408451 15.33333333 6.6      15.31147541
 17.78333333 22.6      38.      22.23529412 70.      8.5
 18.      7.26086957 44.      15.33333333 17.7972973
 40.      10.42105263 3.5      50.      5.86477273 12.1402439
 70.      15.54545455 35.      16.      84.      15.31147541
 6.75      25.      40.      14.66666667 35.      15.31147541
 31.66666667 10.      14.8      4.      52.      5.86477273
 44.      60.      15.31147541 15.04347826 7.26086957 7.
 17.7972973 15.04347826 12.1402439 20.      6.6      15.54545455
 13.75      11.      11.31944444 15.31147541 31.25      15.31147541
 45.      7.95238095 5.86477273 17.      22.23529412 15.01408451
 17.7972973 13.      11.31944444 11.5      15.33333333 35.
 17.      2.25      25.      11.5      6.25      7.26086957
 7.26086957 24.      16.      40.      10.625      12.
 22.23333333 26.27777778 44.      45.      24.5      23.
 50.      60.      10.75862069 40.      20.      45.
 45.      7.91666667 28.      7.91666667 40.16666667 10.
 15.31147541 10.66666667 48.      16.      17.7972973 16.33333333
 40.      3.      17.78333333 5.86477273 22.      30.
 10.625      5.      4.      18.14      17.7972973
 25.02469136 7.26086957 10.      12.1402439 70.      8.29473684
 7.91666667 20.      15.31147541 7.91666667 22.      25.01408451
 30.90909099 44.      40.      15.33333333 15.01408451 30.
 21.25      22.23529412 15.31147541 5.8      5.86477273 12.1402439
 28.      80.      7.91666667 30.      14.5      10.42105263
 13.      15.33333333 50.07142857 7.60294118 50.      31.91666667
 27.9      19.33333333 4.      20.      10.5      15.66666667
 3.      18.5      12.      10.      31.91666667 40.
 7.91666667 3.      40.16666667 13.      35.      20.
 7.95238095 40.      6.      2.5      3.      3.
 25.02469136 27.7972973 15.31147541 35.      12.1402439 7.95238095
 28.      20.      10.      10.      35.      41.      7.91666667
 15.01408451 20.      14.8      48.      12.1402439 17.78333333
 5.86477273 25.02469136 32.      22.8      14.8      27.66666667
15.01408451 20.      14.8      48.      12.1402439 17.78333333
8.6477273 18.      24.      7.91666667 15.31147541 17.78333333
29.      5.97058824 15.01408451 25.02469136 22.8      25.
31.66666667 15.01408451 4.18181818 7.91666667 35.      7.95238095
10.      19.33333333 11.5      30.      30.90909099 35.
7.26086957 10.      9.2      15.01408451 20.      30.
44.      13.      15.31147541 27.06666667 9.5      10.75862069
7.91666667 4.      2.5      4.83333333 26.27777778 90.
7.95238095 26.66666667 27.33333333 6.05405405 8.07741935 12.1402439
14.66666667 10.      11.      15.01408451 23.14285714 8.07741935
8.89473684 5.      23.14285714 30.      22.      10.625
9.0047619 23.      16.33333333 5.86477273 11.31944444 44.
10.      4.18181818 33.75      25.      11.31944444 32.5
5.86477273 25.02469136 32.      22.8      14.8      27.66666667
15.01408451 20.      14.8      48.      12.1402439 17.78333333
8.6477273 18.      24.      7.91666667 15.31147541 17.78333333
50.07142857 15.      6.05405405 25.      13.      7.91666667
18.      44.      25.02469136 10.66666667 40.16666667 8.5
30.90909099 32.5      4.05      12.      15.      7.
7.87837037 5.9047619 ]

In [134]: import numpy as np
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import Transformer
import xgboost as xgb

# Preprocess the features
numeric_features = ['min_price', 'max_price']
categorical_features = ['commodity']

# Scale the numeric features
numeric_transformer = StandardScaler()

# One-hot encode the categorical features
categorical_transformer = OneHotEncoder(sparse=False)

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ])

# Fit the preprocessor on the entire data
preprocessor.fit(X)

# Transform the testing data (including the first row)
X_test_preprocessed = preprocessor.transform(X_test)

# Extract the first row from the preprocessed testing data
selected_row = X_test_preprocessed[0]

# Reshape the row to match the model's input shape
selected_row_reshaped = np.reshape(selected_row, (1, -1))

# Define the XGBoost regressor
model = xgb.XGBRegressor()

# Train the XGBoost regressor on the entire data
model.fit(preprocessor.transform(X), Y)

# Make the prediction using the first row
prediction = model.predict(selected_row_reshaped)

# Display the prediction
print("Predicted modal price:", prediction)

Predicted modal price: [25.977482]

In [135]: X_test

Out [135]:
```

	commodity	min_price	max_price
1760	1	0.275397	0.506581
3326	0	0.161962	1.536581
1176	1	1.283566	1.349308
3176	0	-0.665561	-0.835540
2099	0	2.291736	2.285672
...
2510	0	-0.262293	-0.710692
1762	0	-0.060659	-0.523419
2859	1	-0.665561	-0.991601
423	1	-0.665561	-0.835540
2990	0	-0.732772	-0.523419

800 rows × 3 columns

```
In [ ]:
```